The Perks of Being a First Mover

Does White's Advantage in Chess Change with Rating, Opening and Time Control?

CS215 Open Project · Final Report · Group 02

Sanidhya (25b0944), Daksh Maahor (25b0974), Sahil (25b0991), VadaPavJai (25b1030)

Mentor: Namish Bansal

https://github.com/cs215-cse-iitb/dai-aug-2026-openproject-group-2

Section 1: Meta Information

1. Open Project Title	The Perks of Being a First Mover: Does White's Advantage in Chess Change with Rating, Opening and Time Control?
2. Group ID	"02"
3. Mentor Name	Namish Bansal
4. Group Members and Roll Numbers	Sanidhya (25b0944); Daksh Maahor (25b0974); Sahil (25b0991); VadaPavJai (25b1030)
5. Honor Code and Discussions	None
6. LLM Usage Links	**[Readable links to all chat sessions / prompts used. If an AI agent was used, export its report as PDF, host it and link it here.]**
7. Video Presentation	**[link to video]**
8. GitHub Repository	https://github.com/cs215-cse-iitb/dai-aug-2026-openproject-group-2
9. GitHub Handles and Initial Commit	
Sanidhya	https://github.com/SS3642 earliest commit: October 21, 2026
Daksh Maahor	https://github.com/Daksh-Maahor earliest commit: November 3, 2026
Sahil	https://github.com/Sahil-Deo-IITB earliest commit: November 3, 2026
VadaPavJai	https://github.com/nauseated-potato earliest commit: October 18, 2026
10. Contribution Breakdown	Unequal

Work split (for reference).	Sanidhya	Statistical modelling: fitting h, confidence intervals and the z-tests.
	Daksh Maahor	Data cleaning and opening classification; win/draw/loss rates.
	Sahil	Data engineering: the C++ streaming parser and the Colab pipeline
	VadaPavJai	Visualisation: figures, the Plotly dashboard and the report.

Section 2: Technical Content

Core Problem Statement

We measure the first-move advantage on Lichess. Instead of raw win rates, we fit White's advantage as a number of *rating points*, h, after correcting each game for the rating gap between the players. Under the Elo model a player d points above the opponent scores $E(d) = 1/(1 + 10^{-d/400})$; we give White an extra h points, $E_i(h) = 1/(1 + 10^{-(R_i^{\mathrm{W}} + h - R_i^{\mathrm{B}})/400})$, and choose the h that makes expected scores equal actual scores ($S_i = 1, \frac{1}{2}, 0$). This is the maximum-likelihood estimate (Newton's method; standard error from the Fisher information). Overall, $h =$ **[HAll]** $\pm$ **[HAllCI]** rating points. We test three hypotheses:

- **H1 (Rating).** h peaks at intermediate ratings and shrinks at 2200+.
- **H2 (Openings).** Four aggressive 1.e4 gambits (King's, Evans, Danish, Smith-Morra) carry a larger h than four drawish openings (Petrov, Berlin, QGD, French Exchange), and the drawish ones have more draws at every rating.
- **H3 (Time control).** h grows from Bullet to Classical.

Data Source(s)

Lichess Open Database (`https://database.lichess.org`, CC0), standard rated games for **[MonthList]**: one zstd-compressed PGN per month, tens of GB each, openly downloadable. We use seven of 17 header tags: `Result`, `WhiteElo`, `BlackElo`, `Event`, `TimeControl`, `Termination`, `Opening`. We read **[GamesRead]** games and keep **[GamesUsed]** (**[PctUsed]**).

Month	Games read	Games kept	Kept
[filled in by white_advantage]			
Total	**[GamesRead]**	**[GamesUsed]**	**[PctUsed]**

Tools and Software Packages

A C++ program streams the zstd files, reads only header lines, and counts games by opening, speed, 100-point rating bracket and rating gap (two months on two threads, run on Google Colab, streaming directly from Lichess). It writes `counts.csv` (a few MB); all tables, tests, figures, this report (LaTeX with pgfplots) and an interactive Plotly dashboard (`output/final/dashboard.html`) are rebuilt from it in seconds. **[list exact Python/R packages and versions used for the analysis, if any]**

Challenges

Cleaning and design. We drop games with no result (**[SkipNoResult]**), abandoned games (**[SkipAbandoned]**), missing ratings (**[SkipNoRating]**) and rating gaps above 100 (**[SkipGap]**). Games go into 100-point brackets by average rating and into a speed from `Event`. Openings are matched by Lichess name prefix (excluding gambit side-lines inside the drawish openings) and cover **[OpeningGames]** games (**[OpeningShare]**). With millions of games every difference has a tiny p-value, so we report effect sizes with 95% CIs and treat a few rating points as small. **[technical hurdles: e.g. file size/streaming, opening-name matching, memory, runtime]**

Results

Overall

Over all kept games White wins **[WhiteAll]**, draws **[DrawAll]** and loses **[BlackAll]**. White's average score is **[ScoreAll]** against **[ExpectedAll]** predicted by the ratings alone, which is $h =$ **[HAll]** $\pm$ **[HAllCI]** rating points.

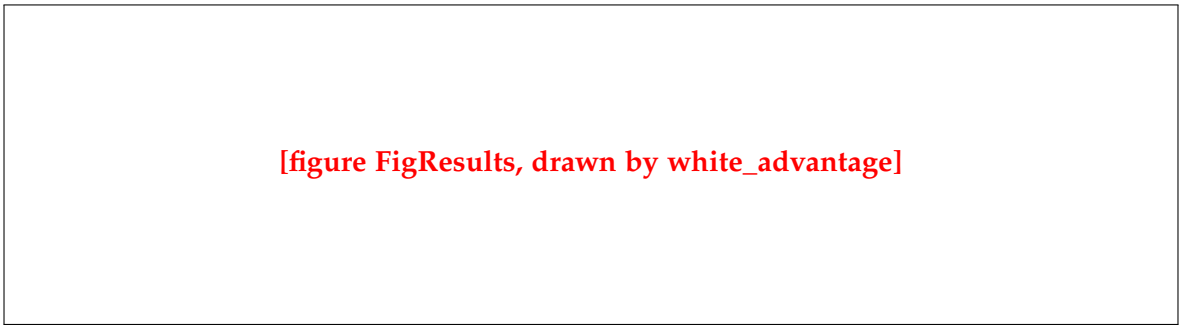


Figure 1: Share of games won by White, drawn and won by Black, by rating (all speeds).

Does the advantage peak at intermediate ratings? (H1)

[interpretation: is H1 supported?]

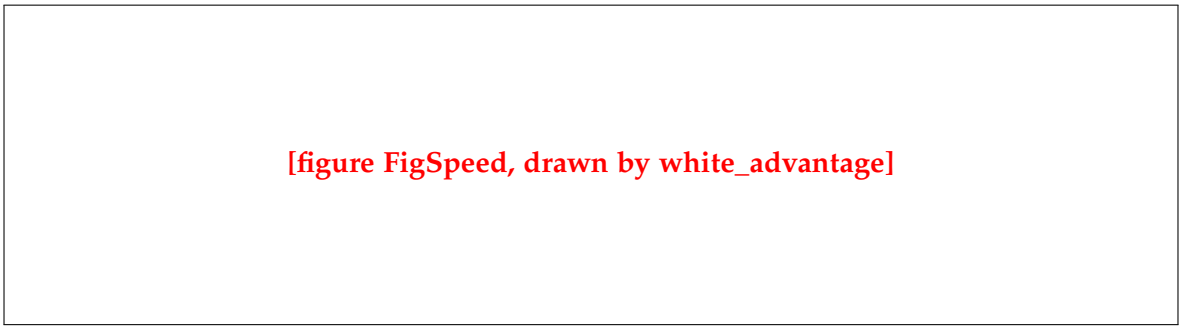


Figure 2: h in each 100-point rating bracket, by speed, with 95% CIs. Brackets with fewer than 1,000 games are left out.

Openings (H2)

Table 1: White's advantage by opening (all speeds). Avg rating = average rating of the games.

Opening	Games	Avg rating	White	Draw	Black	h (pts)
[filled in by white_advantage]						
[filled in by white_advantage]						

Table 2: Aggressive vs drawish openings: h for each group, and the difference with its z-test (rating points, 95% CI), for all speeds and for each of the main speeds.

Speed	Aggressive	Drawish	Difference	z	p
[filled in by white_advantage]					

Over all speeds, h for the aggressive openings minus h for the drawish openings is **[CompareDiff]** rating points ($z =$ **[CompareZ]**, $p =$ **[CompareP]**).

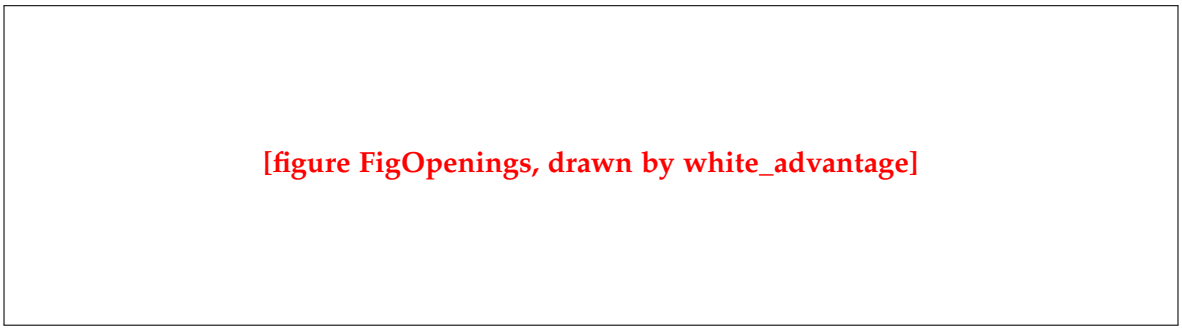


Figure 3: h for each opening (all speeds) with 95% CIs.

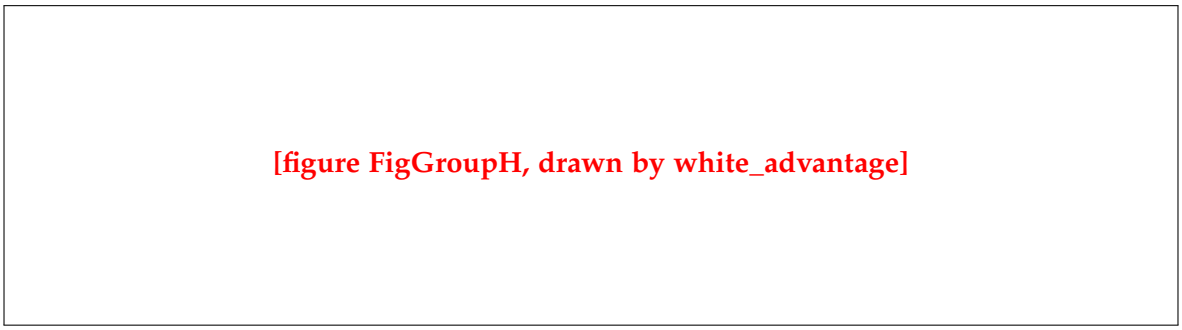


Figure 4: h by rating for the aggressive and the drawish openings, with 95% CIs.

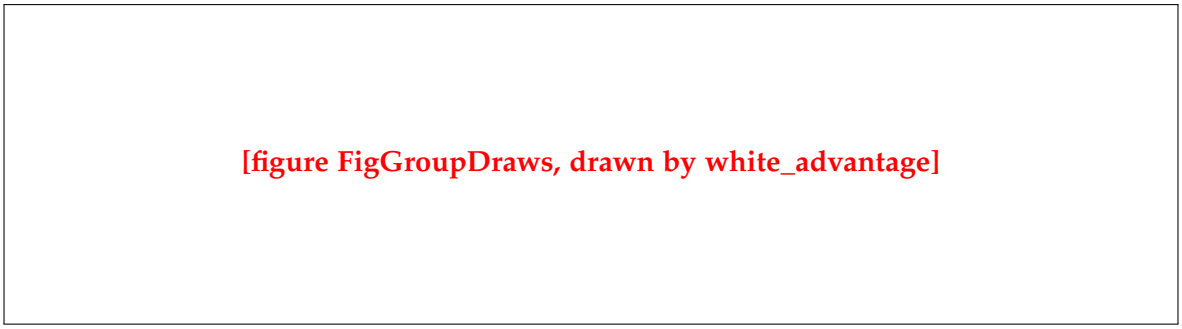


Figure 5: Draw rate by rating for the aggressive and the drawish openings.

[interpretation: is H2 supported, and what does it say about White's advantage?]

Time controls (H3)

Table 3: White's advantage by speed (all openings). White / Draw / Black = share of games won by White, drawn, won by Black. h with 95% CI.

Speed	Games	Avg rating	White	Draw	Black	h (pts)
[filled in by white_advantage]						

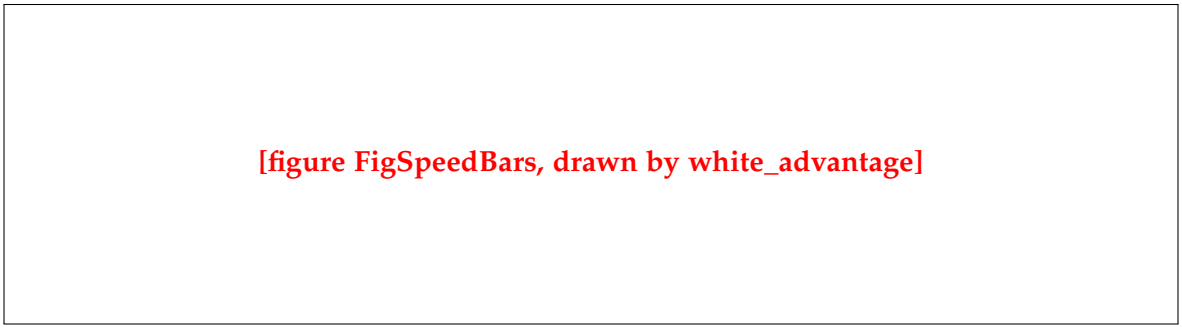


Figure 6: h for each speed (all openings) with 95% CIs.

Table 4: Each speed against the next faster one: h for both, and the difference (slower minus faster) with its z-test (rating points, 95% CI).

Speeds	Slower	Faster	Difference	z	p
[filled in by white_advantage]					

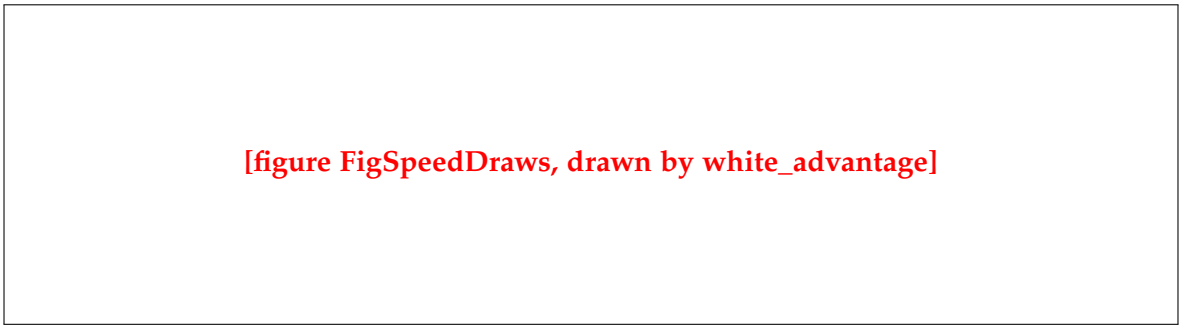


Figure 7: Draw rate by rating for each speed.

[interpretation: is H3 supported?]

Limitations. Opening choice is not random (association, not causation); the two opening groups have different players (see Fig. 4); ratings are noisy and Lichess is not FIDE; games are not independent (active players contribute thousands), so true uncertainty is somewhat larger than our CIs; one site and two months only.

Conclusion

[answer to the three questions in a few sentences]

Section 3: Frequently Answered Questions (FAQ)

Q1. How many games did you analyse and where did they come from?
[GamesUsed] standard rated Lichess games from **[MonthList]** (**[GamesRead]** read, **[PctUsed]** kept), from the CC0 Lichess Open Database.

Q2. What is h and what is its overall value?
h is the number of Elo rating points that, added to White's rating, makes the Elo-expected scores match the observed scores. Overall $h =$ **[HAll]** $\pm$ **[HAllCI]**.

Q3. Which games did you exclude and why?
Games with no result, abandoned games, missing ratings and rating gaps above 100 points: very lopsided games say little about the first move.

Q4. Why rating points instead of raw win rates?
A raw win rate mixes the first move with the rating gap and the draw rate. Rating points correct each game for the gap, so groups with different rating mixes are comparable, and draws count as half a point.

Q5. Why not fit a single slope of h against rating?
The shape is not close to a straight line, so a slope would hide it; we fit h per 100-point bracket and read the shape from the plot.

Q6. With millions of games, isn't everything "significant"?
Yes, so we report effect sizes with 95% CIs and treat differences of a few rating points as small whatever the p-value.

Q7. Which openings were compared?
Aggressive: King's, Evans, Danish and Smith-Morra Gambits. Drawish: Petrov, Berlin, Queen's Gambit Declined, French Exchange (gambit side-lines excluded).

Q8. Does the data support the hypotheses?
[H1, H2, H3 each: supported / partly / not supported, with the key number]

Q9. Does the opening *cause* the difference in h?
Not necessarily. Players choose openings, and styles, preparation and ratings differ between groups, so our results show association only.

Q10. What would you do with more time?
[e.g. more months, cluster-robust CIs by player, adjusting for player ratings, FIDE/OTB data]

Section 4: Reflection Essay

This project taught us to turn a vague question ("is White better?") into something measurable: White's advantage as h rating points, fitted by maximum likelihood with confidence intervals, on millions of real games. We learned to stream huge compressed files with a C++ parser, keep the slow counting step separate from the fast analysis, and rebuild every table and figure from one small file.

Unlike textbook problems, the data were messy and nothing told us what the answer should be. We had to decide what to exclude and defend it. With so many games almost every difference was "significant", so the real question became whether an effect was large enough to matter. We also learned that opening choice is not random, so our results show association, not causation.

The open-ended start felt rudderless, but settling on three testable hypotheses (rating, opening, time control) gave us direction. The work was heavier than a normal assignment because data, modelling and writing all had to fit together, and a late change meant re-running everything. It was also the closest we have come to real data work: building a pipeline, checking the numbers, and explaining them to others.

1. Discussion

[what the results mean for White's advantage, written once the real results are in]

Limitations.

- *Choosing an opening is not random.* Players who play the King's Gambit or the Berlin differ in style and preparation, and the opening is chosen by one player (White for the gambits and the French Exchange, Black for the Petrov, Berlin and QGD). Our h corrects for the rating gap in each game, but the results show association, not that the opening causes the difference.
- *The groups have different players.* The comparison of the two opening groups doesn't adjust for their players' ratings. The "Avg rating" column and Figure 4 (which compares the groups bracket by bracket) show how much that matters.
- *Ratings are noisy*, especially for new accounts, which blurs the correction for the rating gap. Lichess ratings also differ from FIDE ratings, so "2200 on Lichess" is not a master.
- *Games are not independent.* Active players contribute thousands of games, so the real uncertainty is somewhat larger than our confidence intervals.
- *One site, two months.* The results describe Lichess players in early 2024.

References

Lichess Open Database, `https://database.lichess.org`. Lichess opening names, `https://github.com/lichess-org/chess-openings`. A. E. Elo, *The Rating of Chessplayers, Past and Present*, Arco, 1978. M. E. Glickman, "Example of the Glicko-2 system", 2012.